

Context Parallelism for Scalable Million-Token Inference

William Arnold

DALgo Inference Reading Group

2025-12-04

Background: Attention (single query)

For a single query vector \vec{q} :

$$\vec{a} = \vec{q}K^T = [\vec{q} \cdot \vec{k}_1 \ \vec{q} \cdot \vec{k}_2 \ \dots \ \vec{q} \cdot \vec{k}_S] = [a_1 \ a_2 \ \dots \ a_S]$$

$$\text{Softmax}(\vec{a}) = \left[\frac{\exp(a_1 - m)}{Z} \ \frac{\exp(a_2 - m)}{Z} \ \dots \ \frac{\exp(a_S - m)}{Z} \right]$$

$$\text{where } m = \max_j a_j, \quad Z = \sum_{j=1}^S \exp(a_j - m)$$

Background: Flash Attention

$$\forall i \in \{1..S\}$$

$$x_i = \vec{q} \cdot \vec{k}_i$$

$$m_i = \max(m_{i-1}, x_i)$$

$$Z_i = Z_{i-1} e^{m_{i-1} - m_i} + e^{a_i - m_i}$$

$$\vec{o}_i = \vec{o}_{i-1} e^{m_i - m_{i-1}} \frac{Z_{i-1}}{Z_i} + \frac{e^{x_i - m_i}}{Z_i} \vec{v}_i$$

At $i = S$, \vec{o}'_S is the correct output for query \vec{q} .¹

Define AttnBlock($\vec{q}, K, V, \vec{o}_{i-1}, m_{i-1}, Z_{i-1}$) \rightarrow (\vec{o}', m, Z):

¹Flash Attention Explained

Ring Attention

Compute attention on pieces of the sequence!

N ranks, give each rank $1/N$ of the sequence

$$\{Q_1, \dots, Q_N\}, \{K_1, \dots, K_N\}, \{V_1, \dots, V_N\}$$

On rank i , keep Q_i and compute $\forall i \in \{1..N\}$

$$\vec{o}_i, m_i, Z_i \leftarrow \text{AttnBlock}(Q_i, \mathbf{K}_i, \mathbf{V}_i, \vec{o}_i, m_{i-1}, Z_{i-1})^{[1]}$$

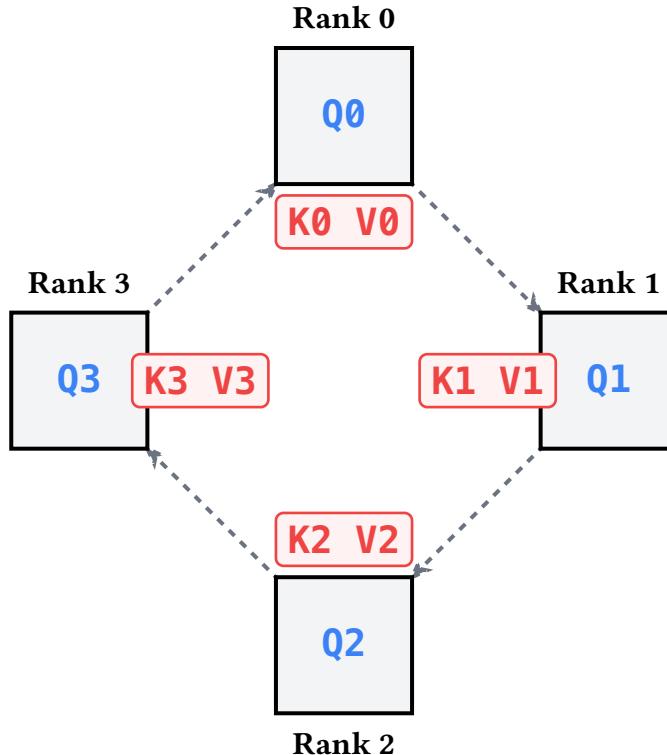
\vec{o} will have the correct output for Q_i

How to get \mathbf{K}_i and \mathbf{V}_i ?

¹ o_i is initialized to zero

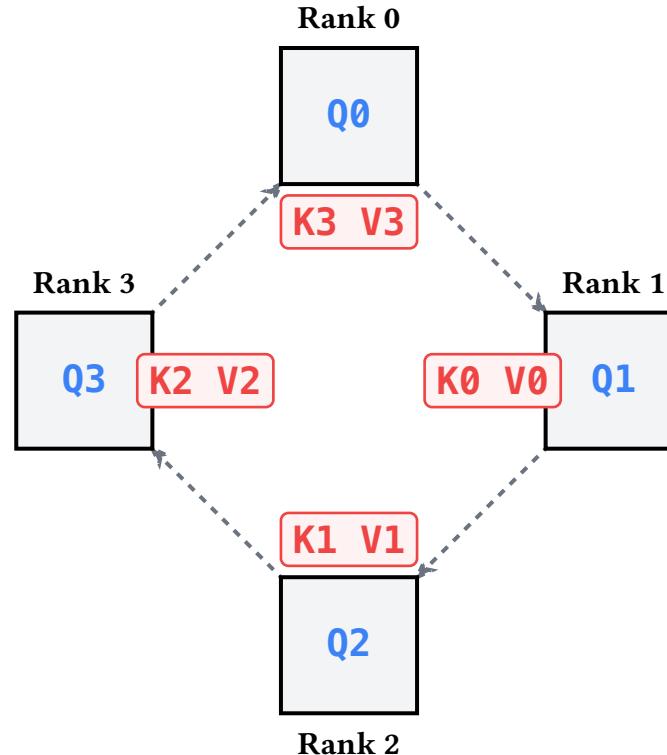
Ring Attention: Communication

Step 1: Each rank computes local attention



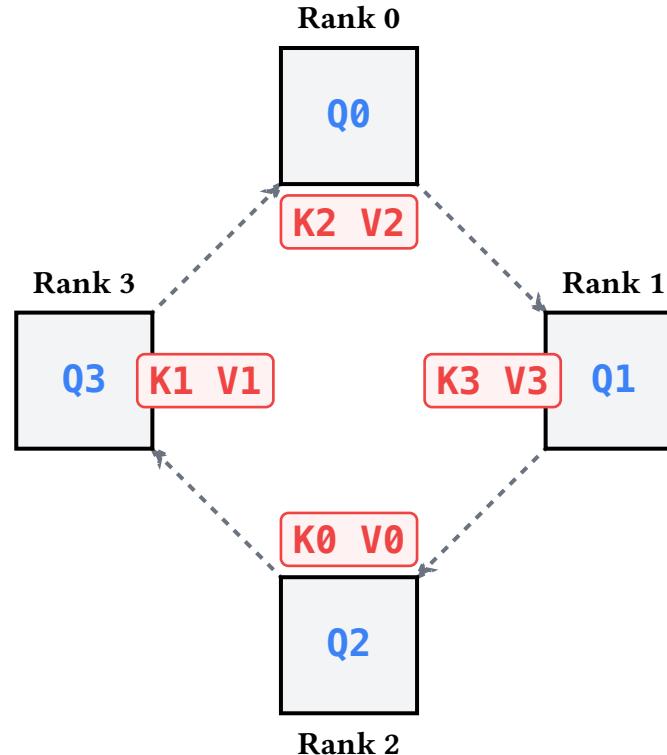
Ring Attention: Communication

Step 2: Send KV to next rank



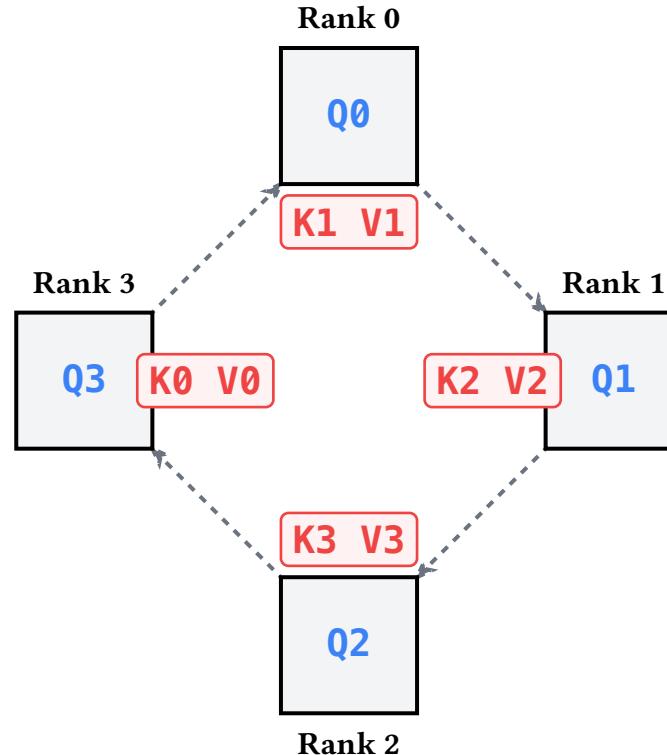
Ring Attention: Communication

Step 3: Send KV to next rank



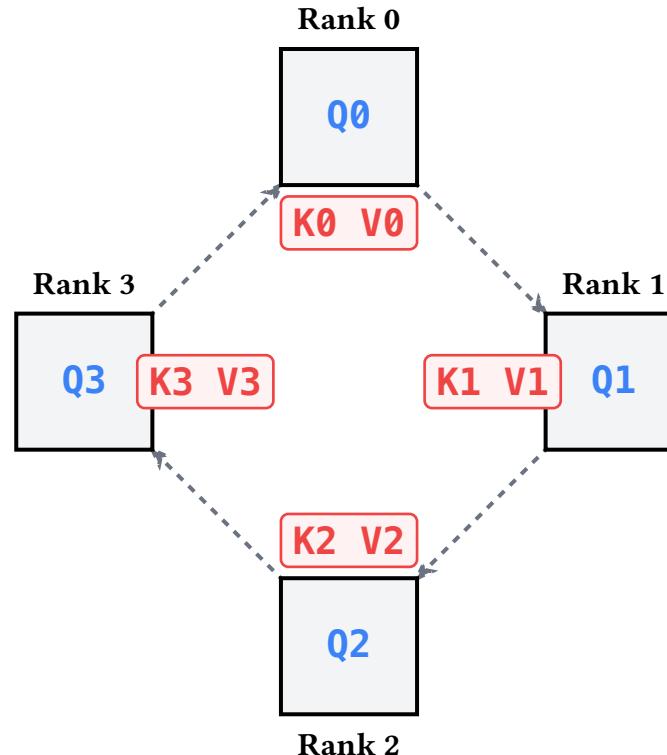
Ring Attention: Communication

Step 4: Send KV to next rank



Ring Attention: Communication

Complete! Each Q has seen all KVs



Ring Attention: Complexity (Pass-KV)

For N ranks, sequence length T , model dim D_q , N_H query heads, N_{KV} KV heads:, e bytes/element

FLOPS	$2T^2 D$
Q bytes	TDe
KV bytes	$2TD\frac{N_{KV}}{N_H}e$

Computation: $\frac{2T^2 D}{N}$ FLOPs

Communication (unidirectional): $2TD\frac{N_{KV}}{N_H}e$ bytes

Ring Attention: Overlap Condition

$$\frac{2T^2 \frac{D}{N}}{C} \geq \frac{2TD\left(\frac{N_{KV}}{N_H}\right)e}{\text{BW}}$$

$$\frac{T}{CN} \geq \frac{N_{KV}e}{N_H \text{ BW}}$$

$$T \geq N\left(\frac{N_{KV}}{N_H}\right)\left(\frac{Ce}{\text{BW}}\right)$$

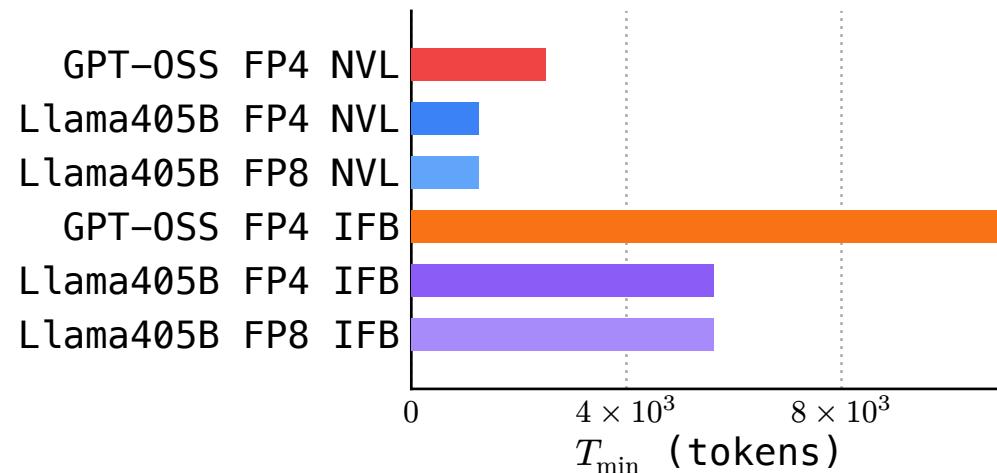
$\frac{Ce}{\text{BW}}$ for Blackwell IFB is ≈ 22500

$\frac{Ce}{\text{BW}}$ for Blackwell NVLink is ≈ 5000

$\frac{Ce}{\text{BW}}$ for Blackwell HMB is ≈ 562.5

When Can We Hide Communication?

Minimum T for communication overlap ($8 \times$ B200)¹:



$$\text{GPT-OSS } \frac{N_H}{N_{KV}} = \frac{64}{8}$$

$$\text{Llama405B } \frac{N_H}{N_{KV}} = \frac{128}{8}$$

¹Infiniband unidirectional @ 200GB/s, 4.5e12 FP8 FLOPS, 9e12 FP4 FLOPS

Ring Attention with Prefixes

With P cached tokens,

FLOPS	$2T(P + T)D$
Q bytes	TDe
KV bytes	$2(P + T)D\frac{N_{KV}}{N_H}e$

Computation: $\frac{2(P + T)TD}{N}$ FLOPs

Communication (unidirectional): $2(P + T)D\frac{N_{KV}}{N_H}e$ bytes

Ring Attention with Prefixes: Overlap Condition

$$\frac{2T(P + T)D}{CN} \geq \frac{2(P + T)D \frac{N_{KV}}{N_H} e}{\text{BW}}$$

$$\frac{T}{CN} \geq \frac{N_{KV}}{N_H} \frac{e}{\text{BW}}$$

$$T \geq N \frac{N_{KV}}{N_H} \frac{Ce}{\text{BW}}$$

! it's the same!

But the new T is *only new tokens*!

Must be 4k+ to hide communication!

This Paper: Context Parallelism over Queries

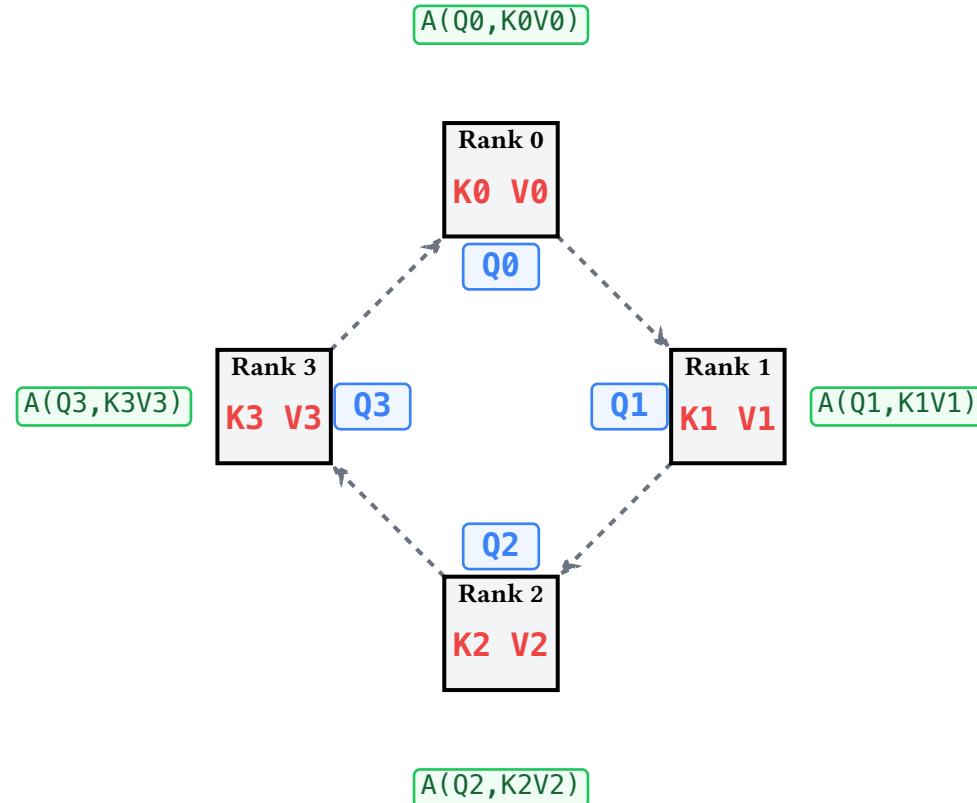
From Meta using 128xH100 with IFB and TCP.

For small T , queries are much smaller than KV!

What if we ring-pass queries?

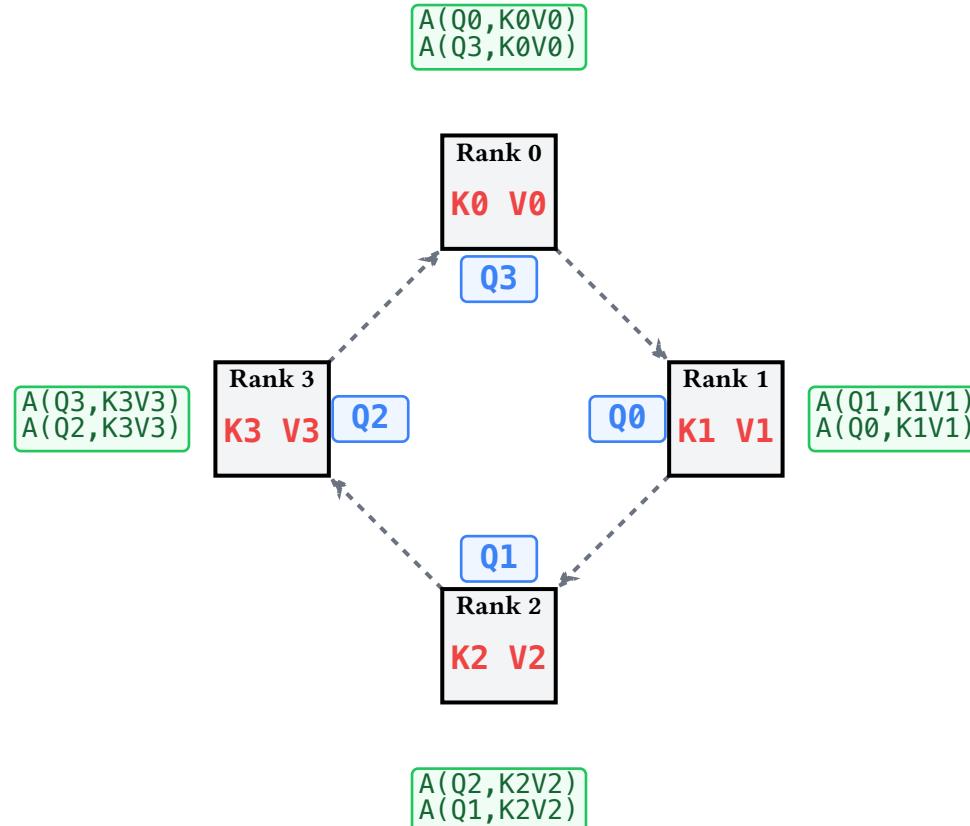
Ring Attention: Pass-Q

Step 1: Each rank computes local attention, stores (o, m, Z)



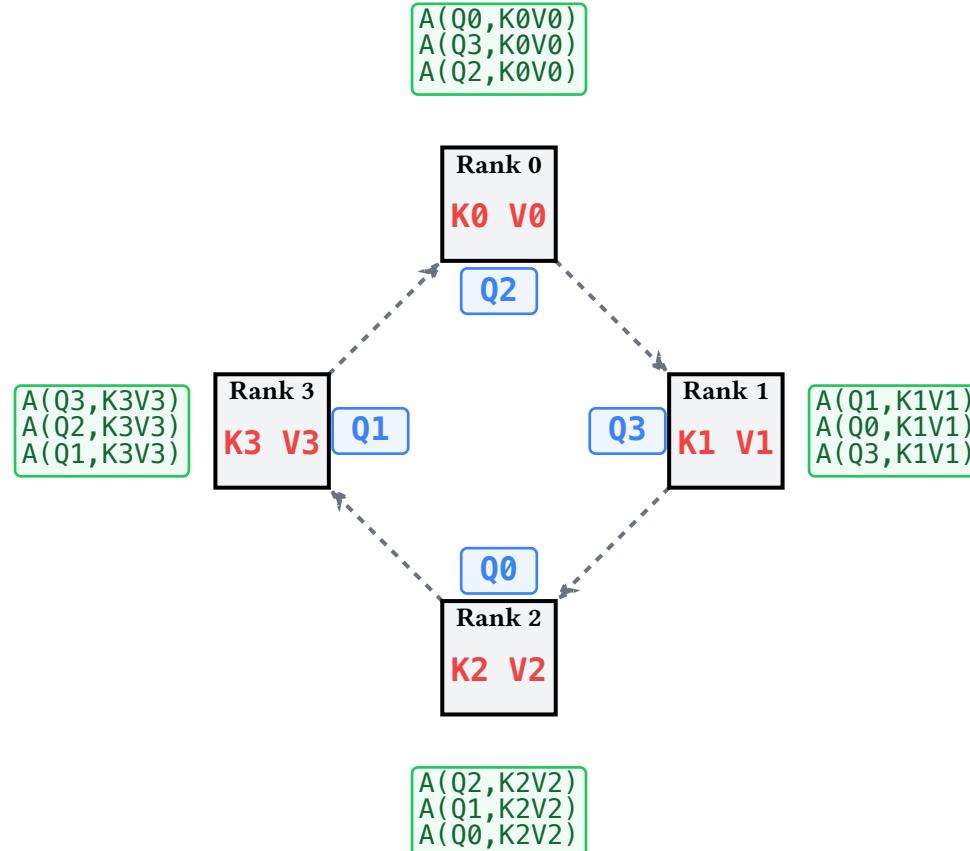
Ring Attention: Pass-Q

Step 2: Send Q to next rank, compute & store



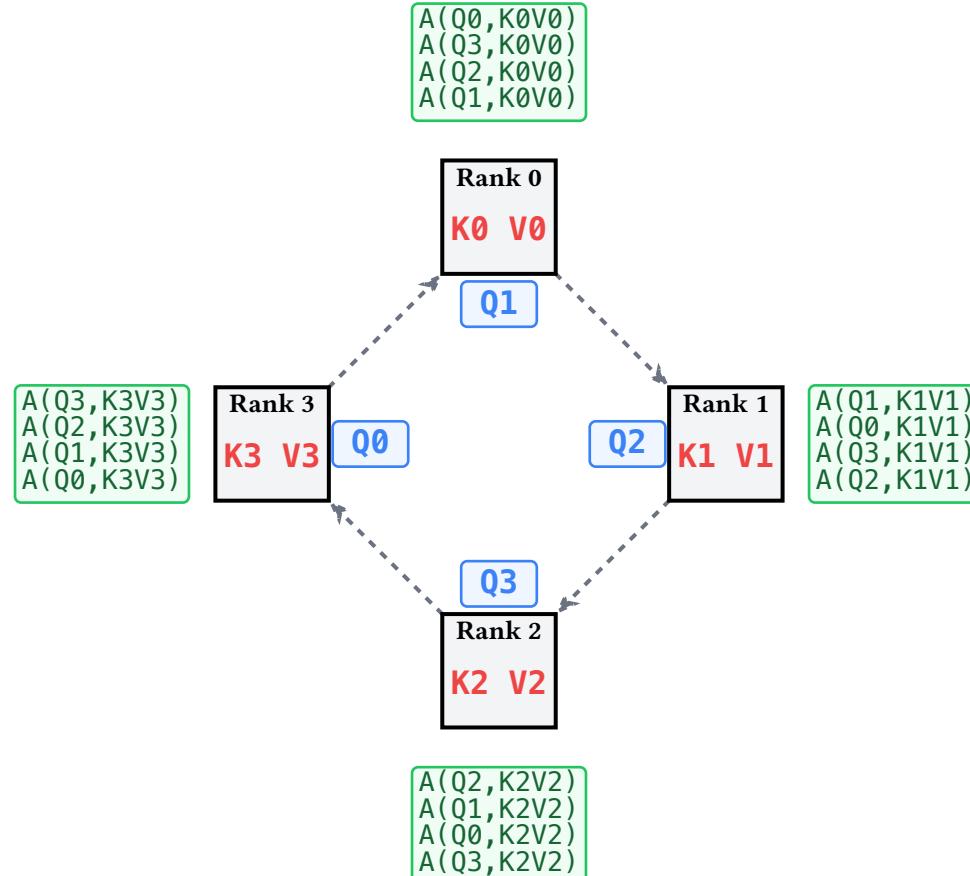
Ring Attention: Pass-Q

Step 3: Send Q to next rank, compute & store



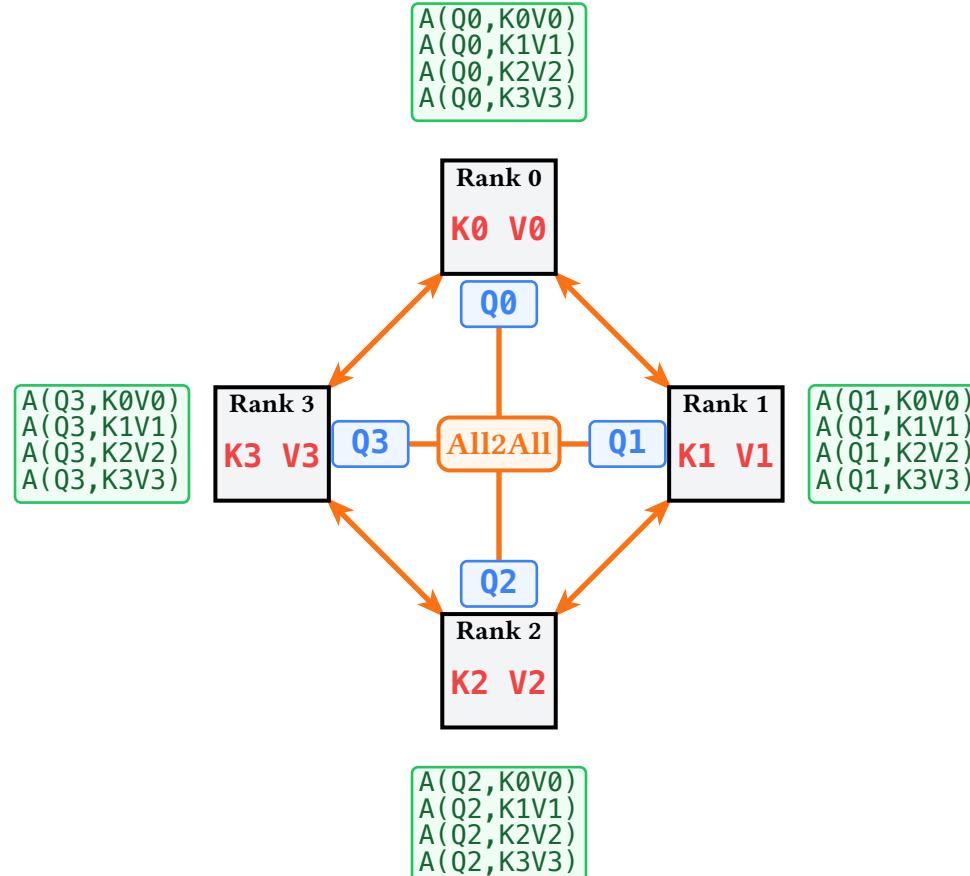
Ring Attention: Pass-Q

Step 4: Send Q to next rank, compute & store



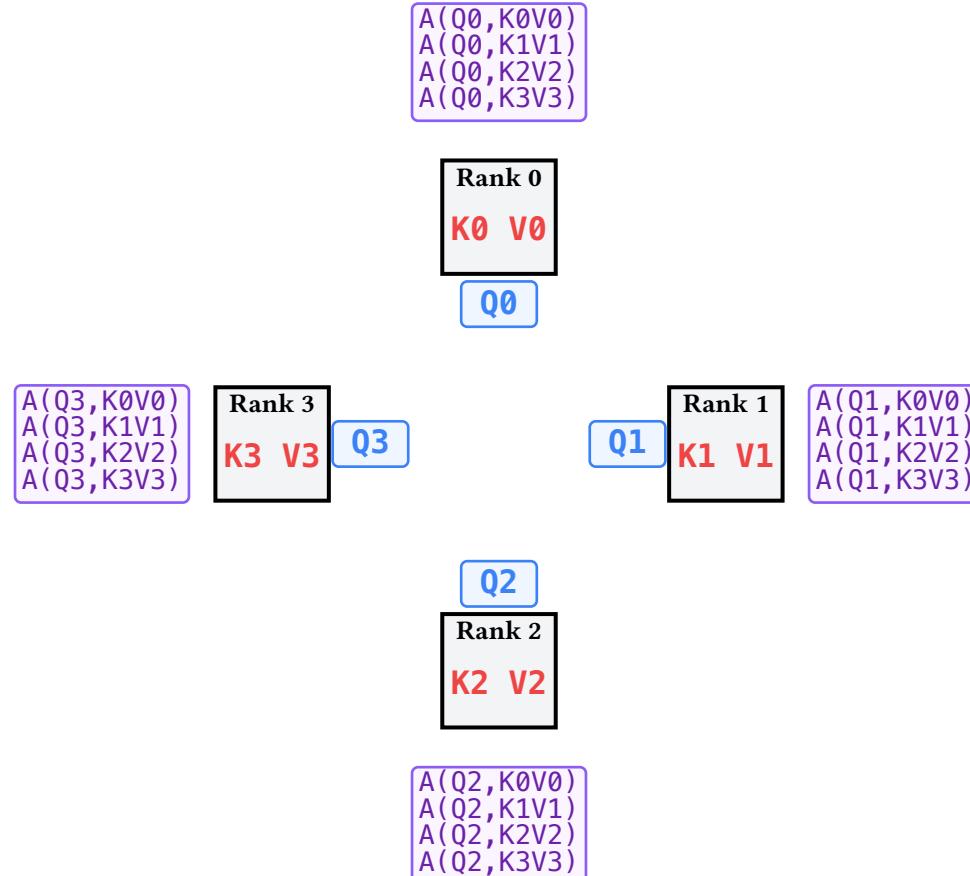
Ring Attention: Pass-Q

All2All: Exchange partial outputs...



Ring Attention: Pass-Q

Done! Each rank has all partials for its own Q



Q-Passing Roofline

FLOPS	$2T(P + T)D$
Q bytes	TDe
KV bytes	$2(P + T)D \frac{N_{KV}}{N_H} e$

Computation: $\frac{2T(P + T)D}{N}$ FLOPs

Communication (unidirectional): TDe bytes

Q-Passing Roofline

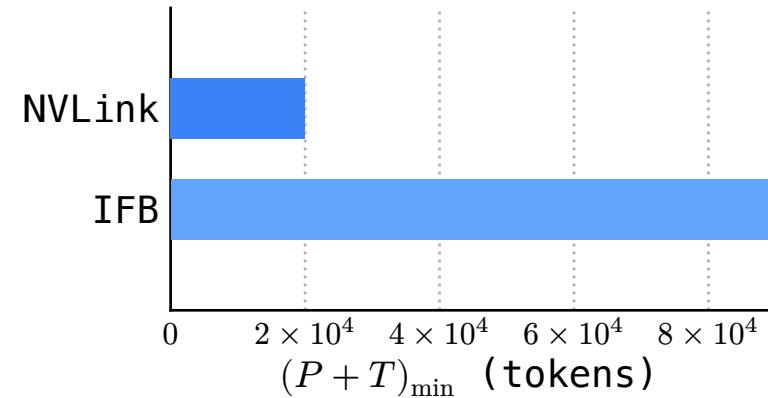
Overlap condition:

$$\frac{2T(P + T)D}{CN} \geq \frac{TDe}{\text{BW}}$$
$$P + T \geq \frac{N}{2} \frac{Ce}{\text{BW}}$$

Doesn't depend on N_H, N_{KV}

Q-Passing Roofline

Minimum $(P + T)$ for communication overlap ($8 \times$ B200):



Note: $C \cdot e$ is the same for FP4 and FP8, so model doesn't matter!

Requires balanced KVs across ranks (round-robin during decode)

All-to-All cost

All-to-All time in a ring is roughly¹

$$T_{\text{all2all}} = \frac{TDe}{4 \text{ BW}}$$

All-to-All time in NVL72 is just $\frac{TDe}{N \text{ BW}}$

Need to check if $T_{\text{all2all}} < (T_{\text{kv,comm}} - T_{\text{kv,compute}})$

¹See this derivation

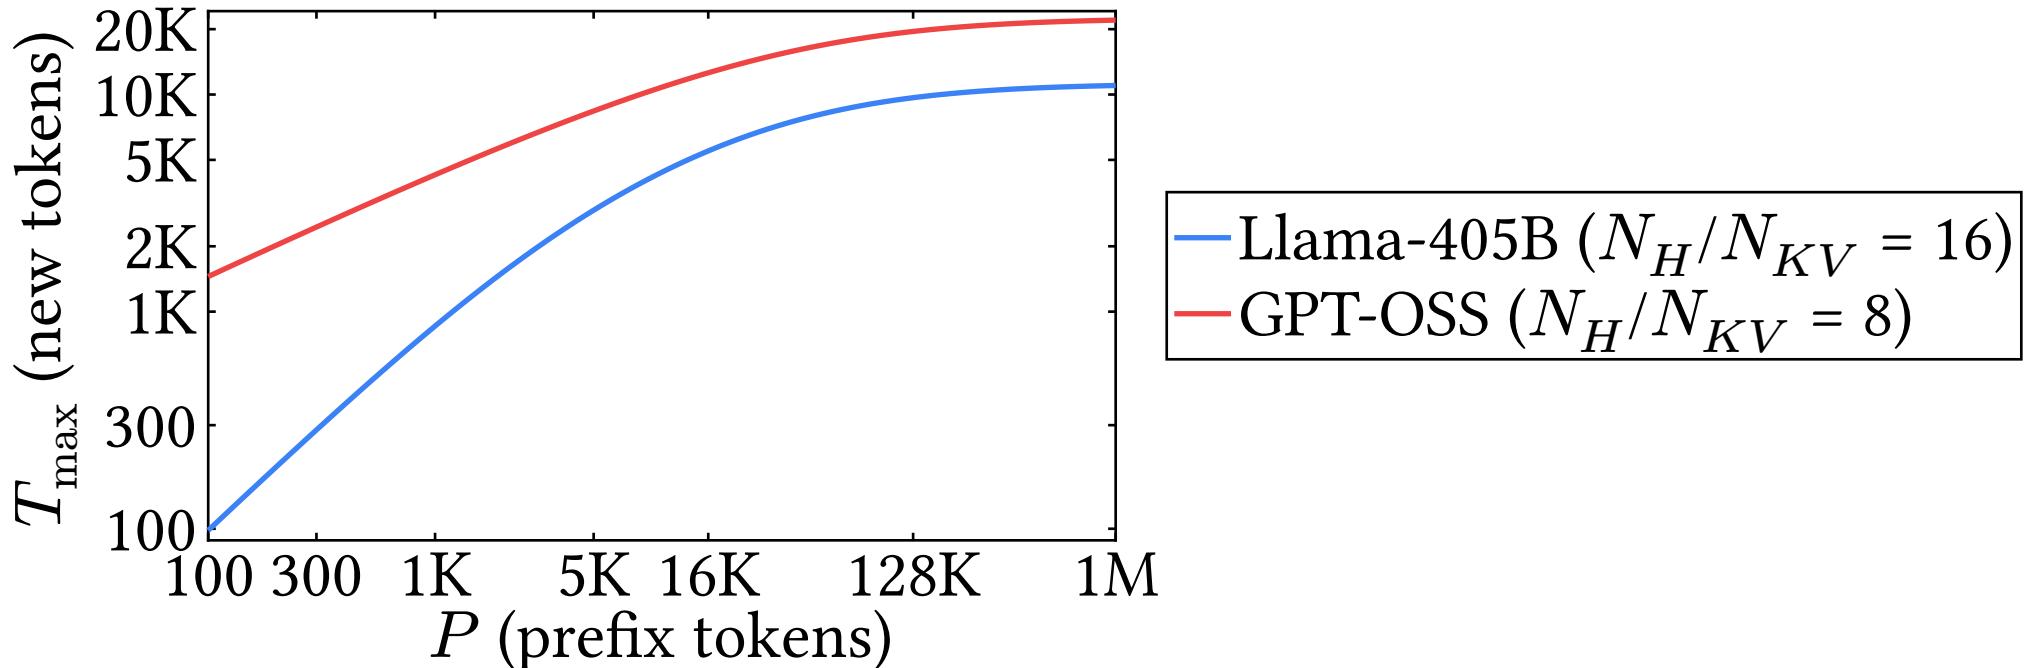
All-to-All cost

$$\begin{aligned} T_{\text{kv,comm}} - T_{\text{kv,compute}} &= 2(P + T)D \frac{N_{KV}}{N_H} \frac{e}{\text{BW}} - \frac{2(P + T)TD}{NC} \\ &= 2(P + T)D \left(\frac{N_{KV}}{N_H} \frac{e}{\text{BW}} - \frac{T}{NC} \right) \\ \frac{TD e}{4 \text{ BW}} &< T_{\text{kv,comm}} - T_{\text{kv,compute}} \end{aligned}$$

Ends up being quadratic... hand it to `sympy`, solve for T and plot

T_{\max} vs Prefix Length

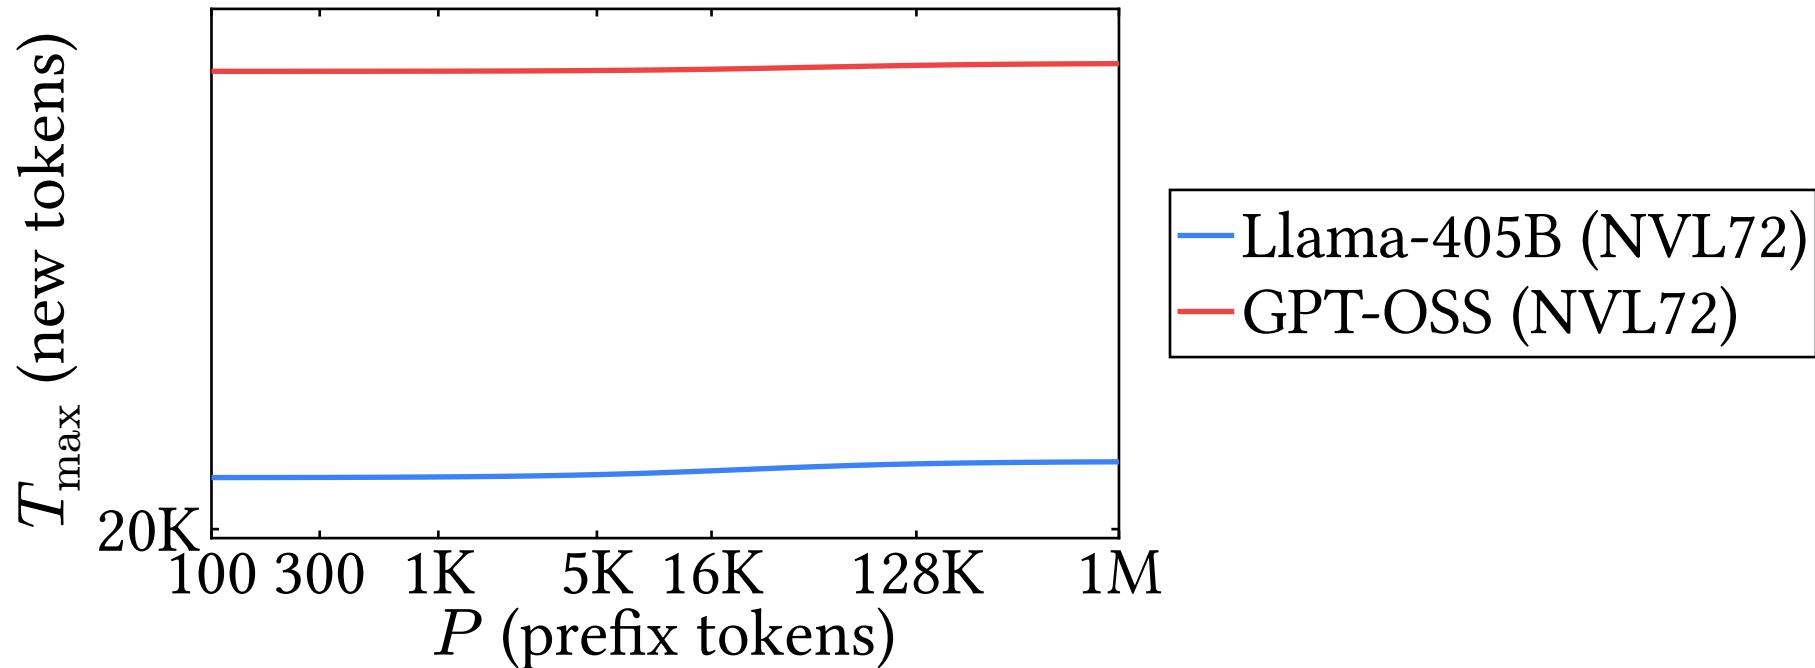
Maximum T where Pass-Q is faster than Pass-KV ($8 \times$ B200 IFB):



Below T_{\max} , use All-to-All (Pass-Q). Above T_{\max} , use Pass-KV.

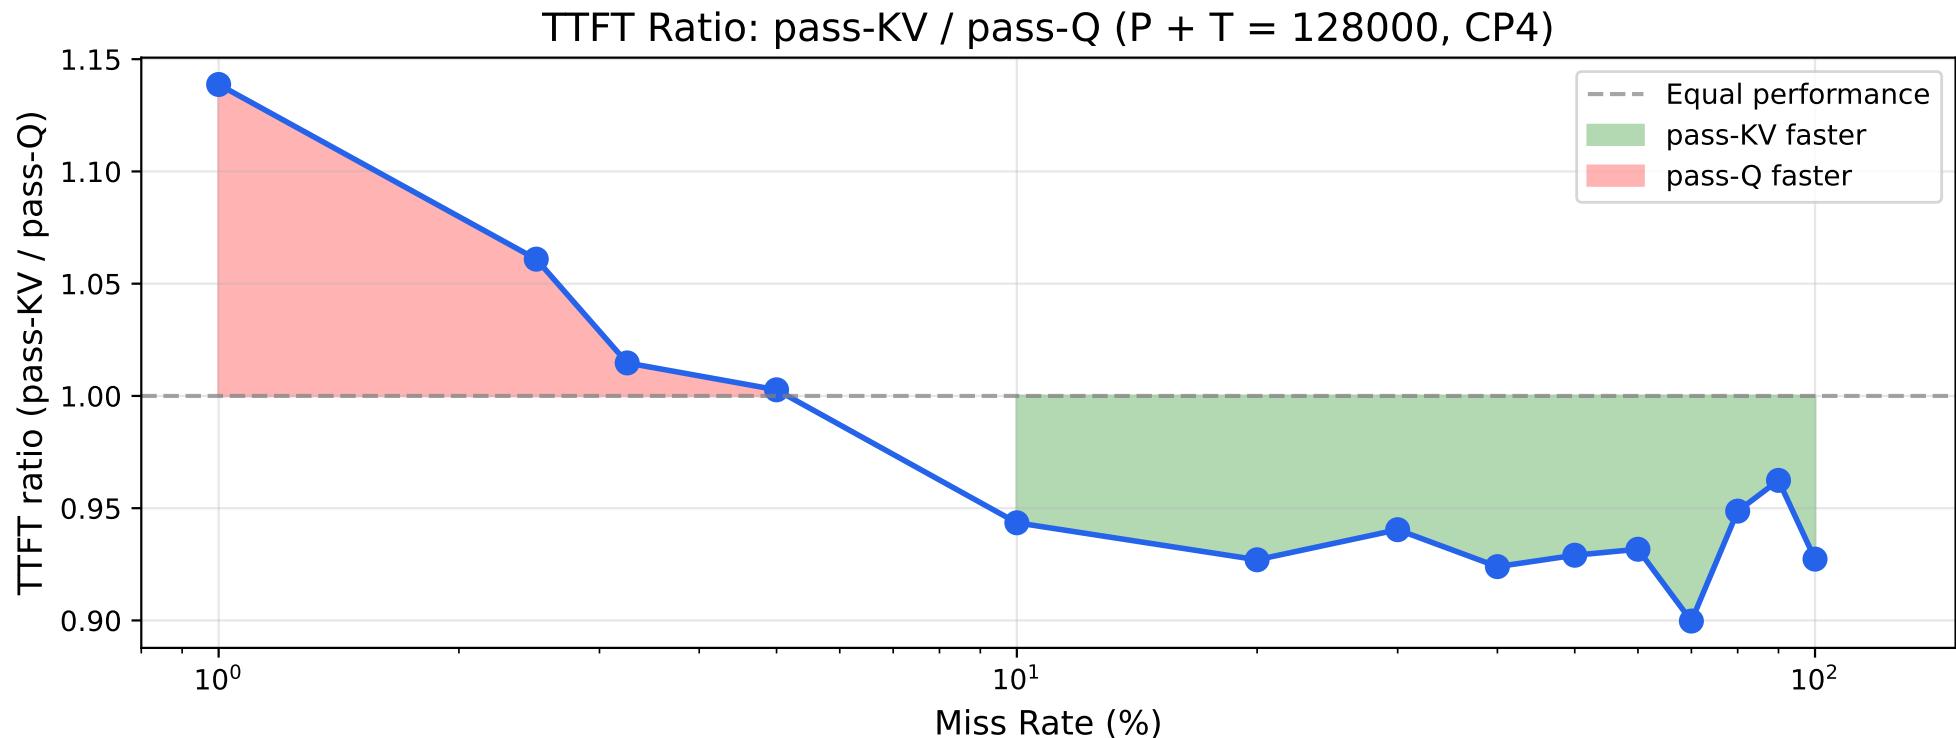
What about NVL72?

All-to-all cost is $\frac{TDe}{N \cdot 4 \cdot \text{BW}}$ and BW is 900GB/s v.s. 200GB/s.



With NVL72's lower all-to-all cost, Pass-Q is optimal up to larger T .

Meta's Hopper Performance

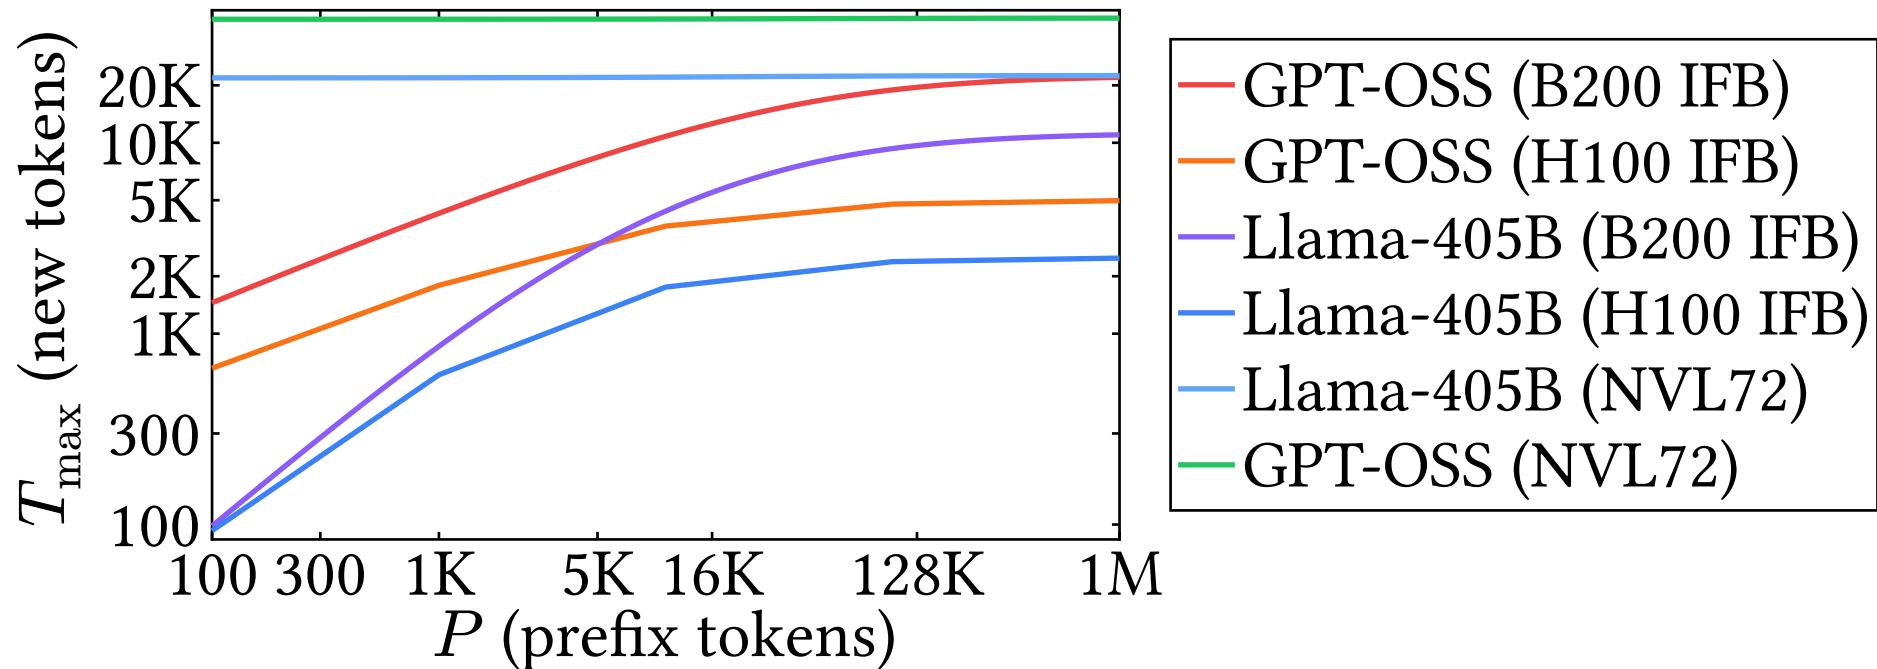


pass-Q is slightly better than pass-KV for very low miss rates

Does the math say the same?

Validating Meta's Hopper Performance

4 x H100 with 200GB/s (unidirectional) IFB



For H100 IFB at $P=128K$, $T_{\max} \approx 2K \approx 2\%$ hit rate! Pretty close!

For NVL72 this works really well!