

# Context Parallelism for Scalable Million-Token Inference

William Arnold

DALgo Inference Reading Group

2025-12-04

# Background: Attention (single query)

For a single query vector  $\vec{q}$ :

$$\vec{a} = \vec{q}K^T = [\vec{q} \cdot \vec{k}_1 \ \vec{q} \cdot \vec{k}_2 \ \dots \ \vec{q} \cdot \vec{k}_S] = [a_1 \ a_2 \ \dots \ a_S]$$

$$\text{Softmax}(\vec{a}) = \left[ \frac{\exp(a_1 - m)}{Z} \ \frac{\exp(a_2 - m)}{Z} \ \dots \ \frac{\exp(a_S - m)}{Z} \right]$$

$$\text{where } m = \max_j a_j, \quad Z = \sum_{j=1}^S \exp(a_j - m)$$

# Background: Flash Attention

$$\forall i \in \{1..S\}$$

$$x_i = \vec{q} \cdot \vec{k}_i$$

$$m_i = \max(m_{i-1}, x_i)$$

$$Z_i = Z_{i-1} e^{m_{i-1} - m_i} + e^{a_i - m_i}$$

$$\vec{o}_i = \vec{o}_{i-1} e^{m_i - m_{i-1}} \frac{Z_{i-1}}{Z_i} + \frac{e^{x_i - m_i}}{Z_i} \vec{v}_i$$

At  $i = S$ ,  $\vec{o}'_S$  is the correct output for query  $\vec{q}$ .<sup>1</sup>

Define AttnBlock( $\vec{q}, K, V, \vec{o}_{i-1}, m_{i-1}, Z_{i-1}$ )  $\rightarrow$  ( $\vec{o}', m, Z$ ):

---

<sup>1</sup>Flash Attention Explained

# Ring Attention

Compute attention on pieces of the sequence!

$N$  ranks, give each rank  $1/N$  of the sequence

$$\{Q_1, \dots, Q_N\}, \{K_1, \dots, K_N\}, \{V_1, \dots, V_N\}$$

On rank  $i$ , keep  $Q_i$  and compute  $\forall i \in \{1..N\}$

$$\vec{o}_i, m_i, Z_i \leftarrow \text{AttnBlock}(Q_i, \mathbf{K}_i, \mathbf{V}_i, \vec{o}_i, m_{i-1}, Z_{i-1})^{[1]}$$

$\vec{o}$  will have the correct output for  $Q_i$

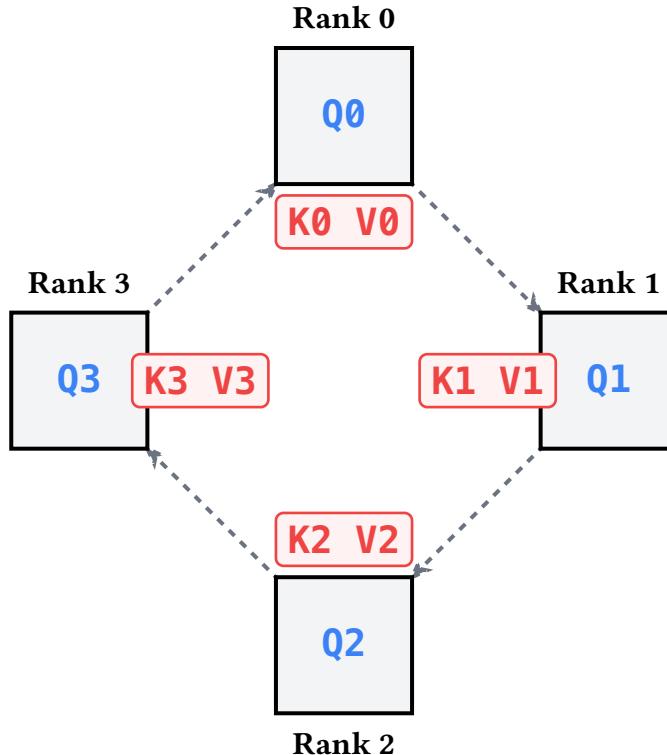
How to get  $\mathbf{K}_i$  and  $\mathbf{V}_i$ ?

---

<sup>1</sup> $o_i$  is initialized to zero

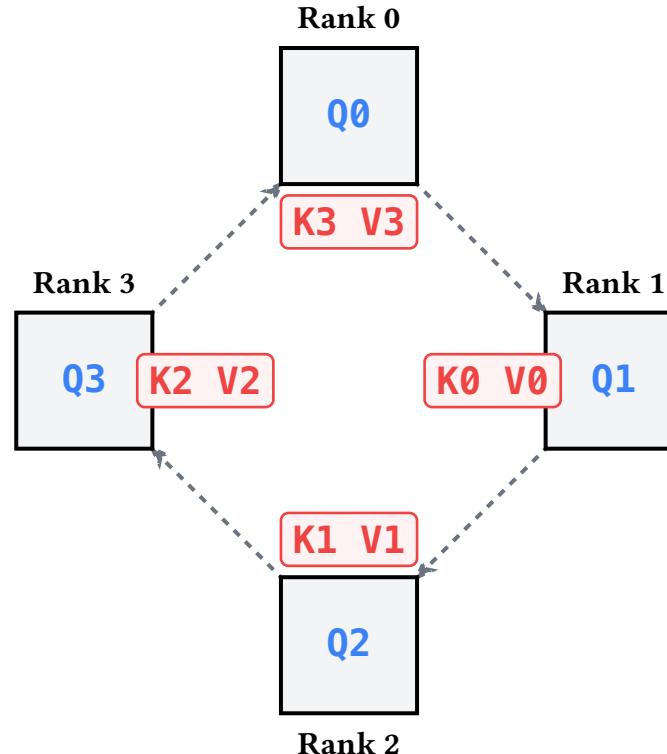
# Ring Attention: Communication

Step 1: Each rank computes local attention



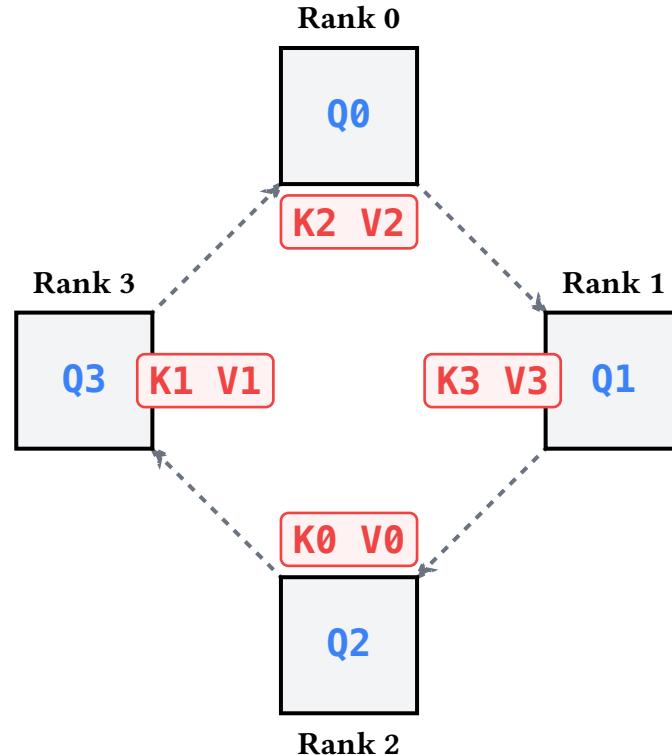
# Ring Attention: Communication

Step 2: Send KV to next rank



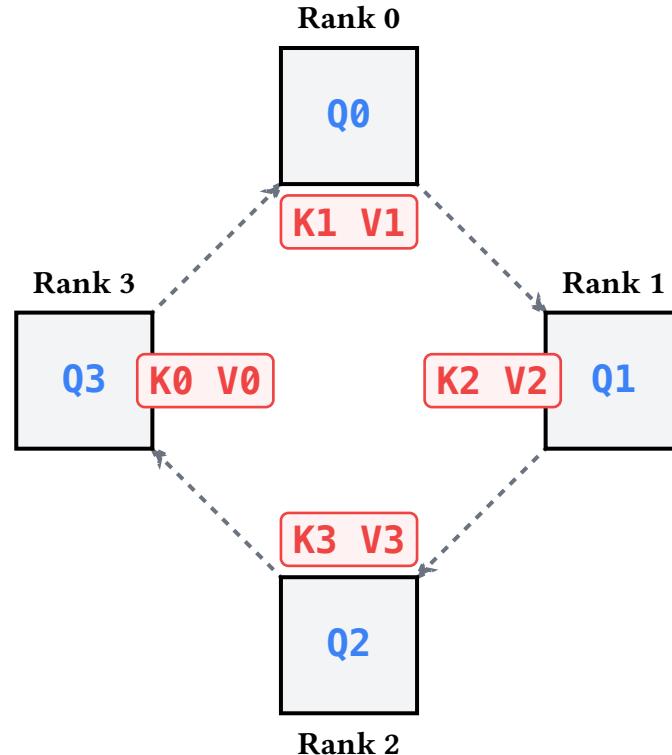
# Ring Attention: Communication

Step 3: Send KV to next rank



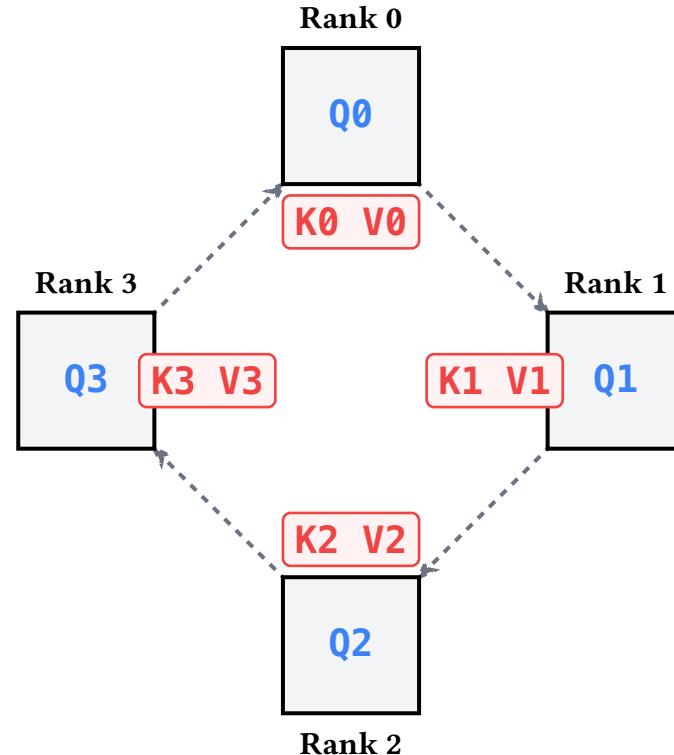
# Ring Attention: Communication

Step 4: Send KV to next rank



# Ring Attention: Communication

Complete! Each Q has seen all KVs



# Ring Attention: Complexity (Pass-KV)

For  $N$  ranks, sequence length  $T$ , model dim  $D_q$ ,  $N_H$  query heads,  $N_{KV}$  KV heads:,  $e$  bytes/element

<b>FLOPS</b>	$2T^2 D$
<b>Q bytes</b>	$TDe$
<b>KV bytes</b>	$2TD\frac{N_{KV}}{N_H}e$

**Computation:**  $\frac{2T^2 D}{N}$  FLOPs

**Communication** (unidirectional):  $2TD\frac{N_{KV}}{N_H}e$  bytes

# Ring Attention: Overlap Condition

$$\frac{2T^2 \frac{D}{N}}{C} \geq \frac{2TD\left(\frac{N_{KV}}{N_H}\right)e}{\text{BW}}$$

$$\frac{T}{CN} \geq \frac{N_{KV}e}{N_H \text{ BW}}$$

$$T \geq N\left(\frac{N_{KV}}{N_H}\right)\left(\frac{Ce}{\text{BW}}\right)$$

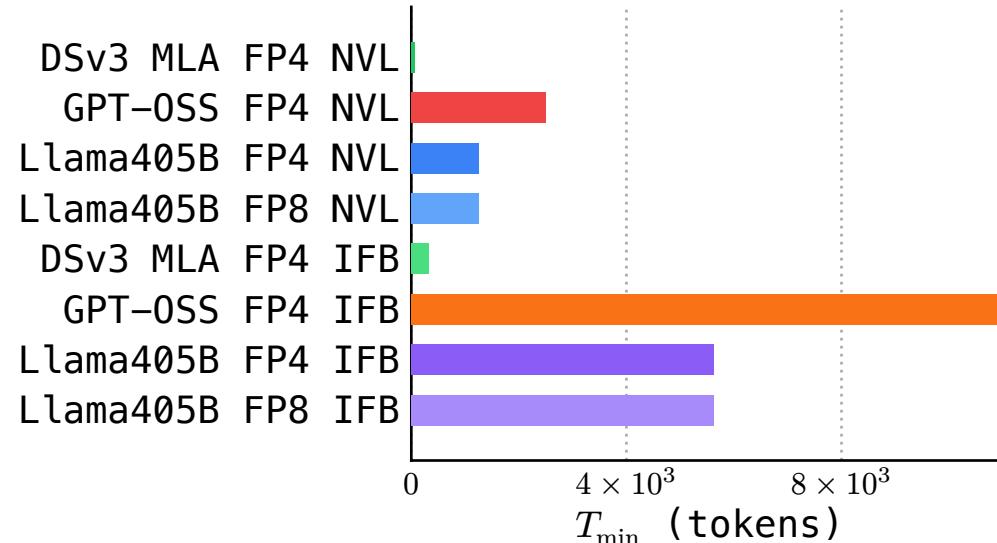
$\frac{Ce}{\text{BW}}$  for Blackwell IFB is  $\approx 22500$

$\frac{Ce}{\text{BW}}$  for Blackwell NVLink is  $\approx 5000$

$\frac{Ce}{\text{BW}}$  for Blackwell HMB is  $\approx 562.5$

# When Can We Hide Communication?

Minimum  $T$  for communication overlap ( $8 \times$  B200)<sup>1</sup>:



Independent of datatype since  $Ce$  is *theoretically* constant. NVL72

<sup>1</sup>Infiniband unidirectional @ 200GB/s, 4.5e12 FP8 FLOPS, 9e12 FP4 FLOPS

# Ring Attention with Prefixes

With  $P$  cached tokens,

<b>FLOPS</b>	$2T(P + T)D$
<b>Q bytes</b>	$TDe$
<b>KV bytes</b>	$2(P + T)D\frac{N_{KV}}{N_H}e$

**Computation:**  $\frac{2(P + T)TD}{N}$  FLOPs

**Communication** (unidirectional):  $2(P + T)D\frac{N_{KV}}{N_H}e$  bytes

# Ring Attention with Prefixes: Overlap Condition

$$\frac{2T(P + T)D}{CN} \geq \frac{2(P + T)D \frac{N_{KV}}{N_H} e}{\text{BW}}$$

$$\frac{T}{CN} \geq \frac{N_{KV}}{N_H} \frac{e}{\text{BW}}$$

$$T \geq N \frac{N_{KV}}{N_H} \frac{Ce}{\text{BW}}$$

! it's the same!

But the new  $T$  is *only new tokens*!

Must be 4k+ to hide communication!

# This Paper: Context Parallelism over Queries

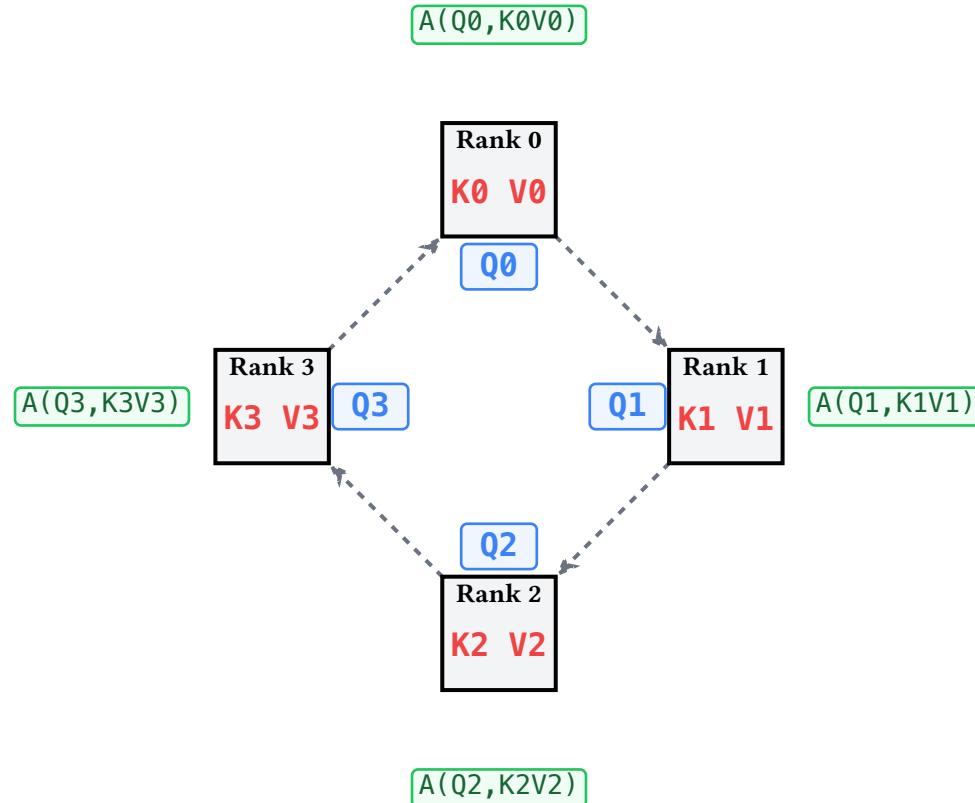
From Meta using 128xH100 with IFB and TCP.

For small  $T$ , queries are much smaller than KV!

What if we ring-pass queries?

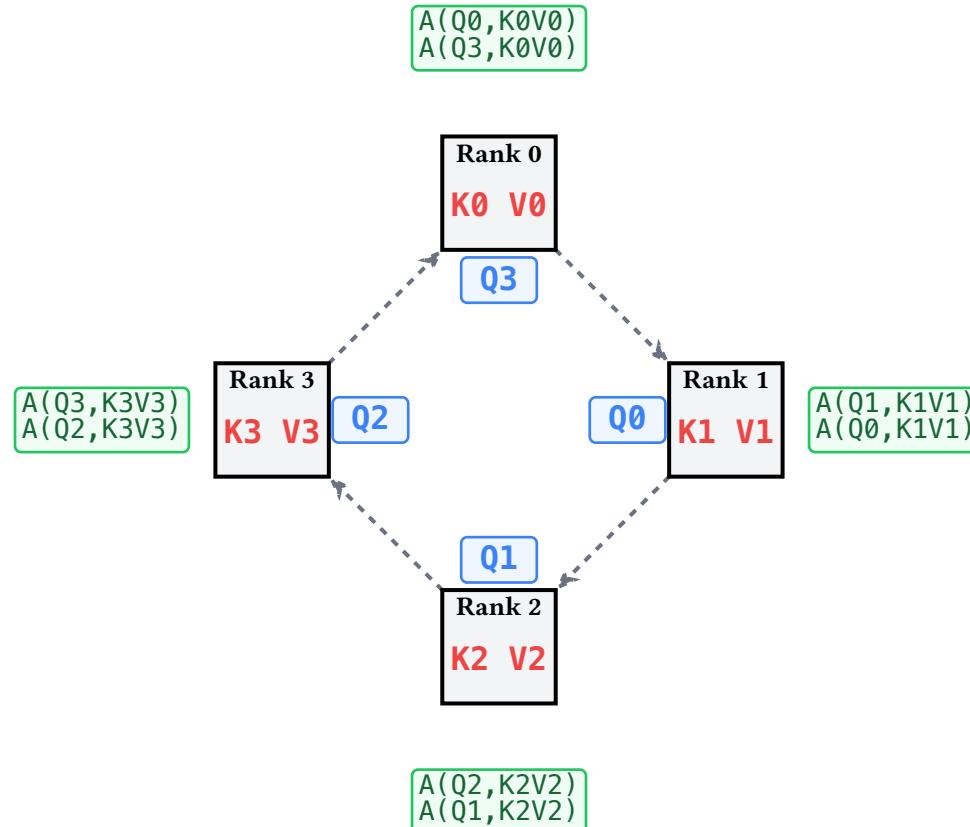
# Ring Attention: Pass-Q

Step 1: Each rank computes local attention, stores  $(o, m, Z)$



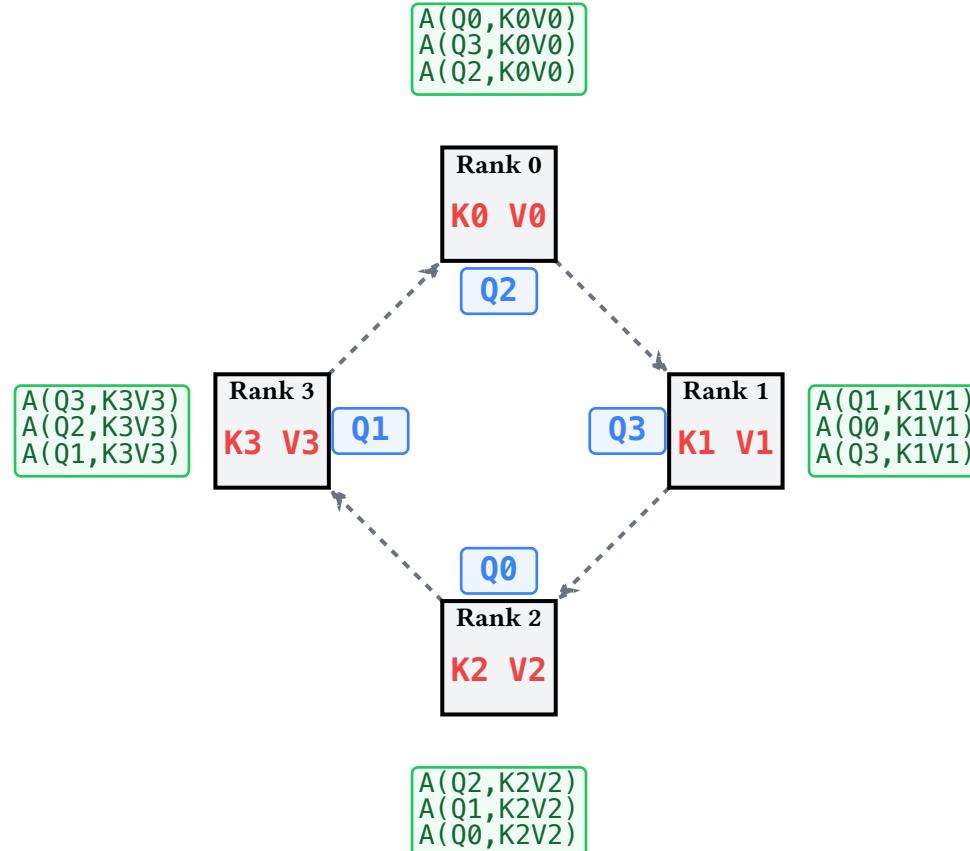
# Ring Attention: Pass-Q

Step 2: Send  $Q$  to next rank, compute & store



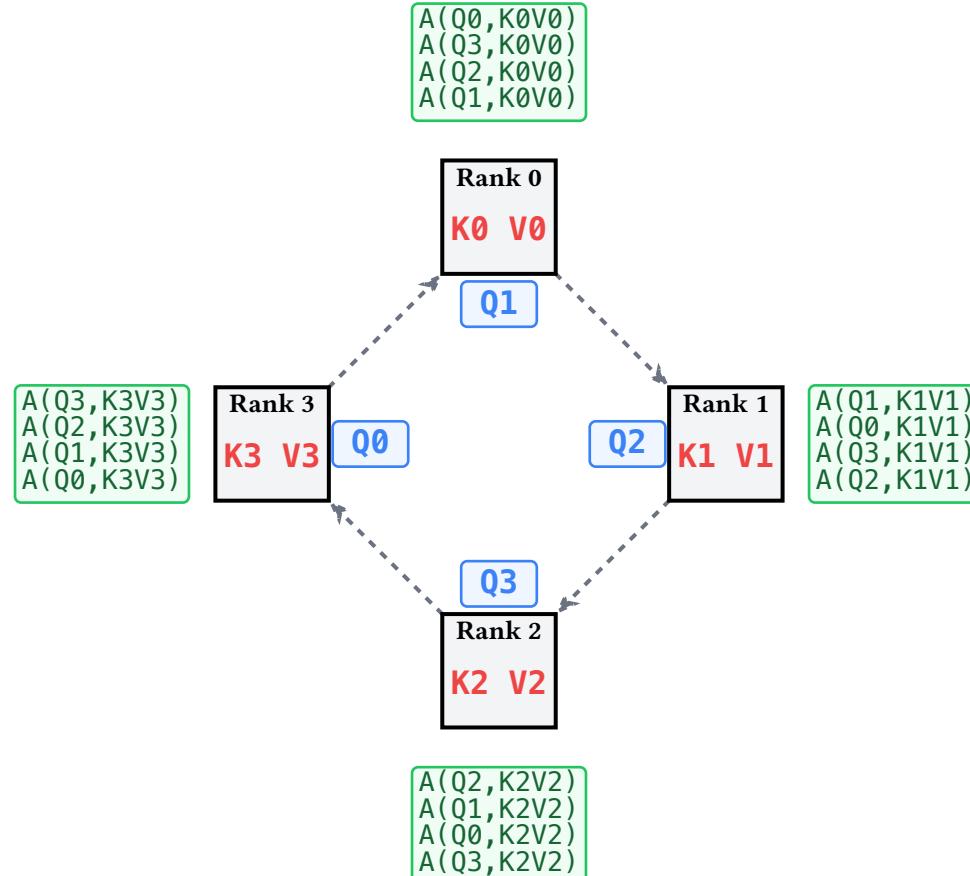
# Ring Attention: Pass-Q

Step 3: Send  $Q$  to next rank, compute & store



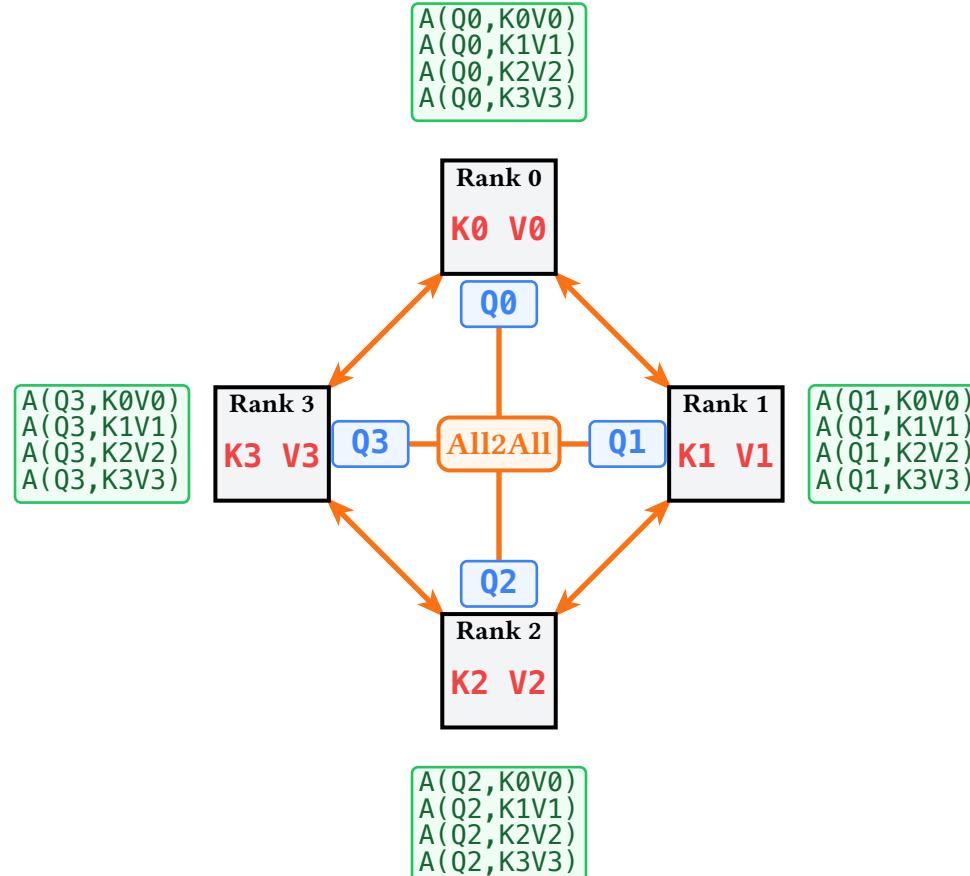
# Ring Attention: Pass-Q

Step 4: Send Q to next rank, compute & store



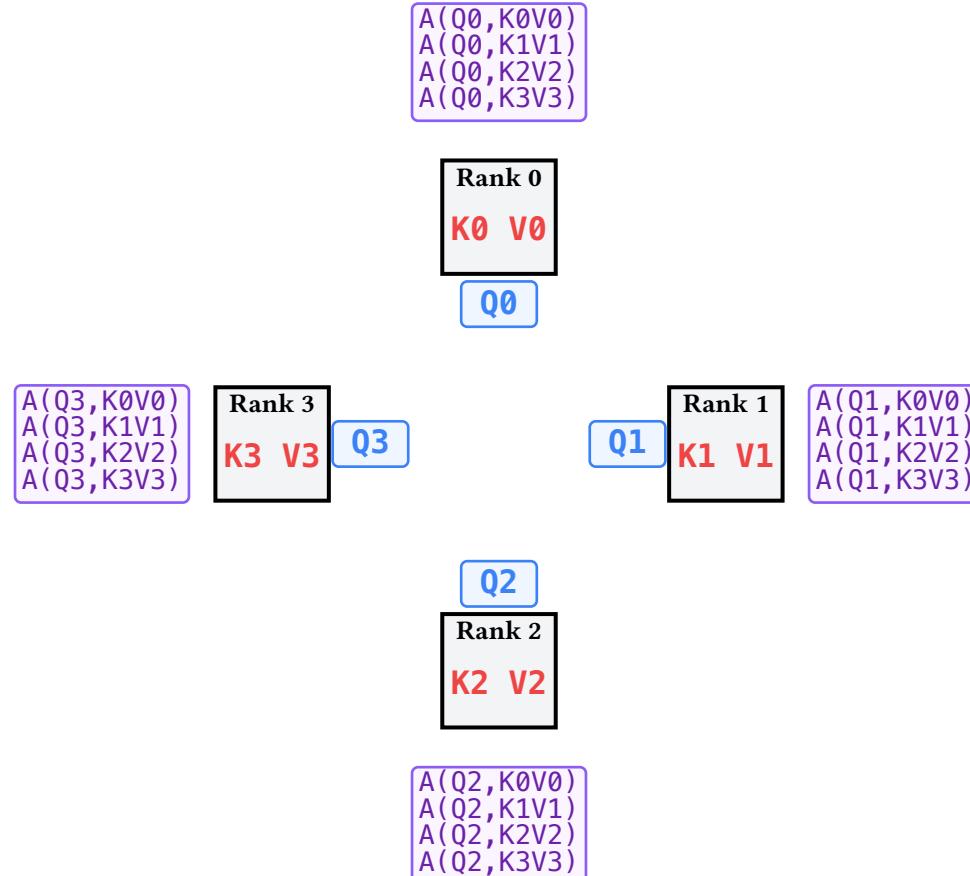
# Ring Attention: Pass-Q

All2All: Exchange partial outputs...



# Ring Attention: Pass-Q

Done! Each rank has all partials for its own Q



# Q-Passing Roofline

FLOPS	$2T(P + T)D$
Q bytes	$TDe$
KV bytes	$2(P + T)D \frac{N_{KV}}{N_H} e$

**Computation:**  $\frac{2T(P + T)D}{N}$  FLOPs

**Communication** (unidirectional):  $TDe$  bytes

# Q-Passing Roofline

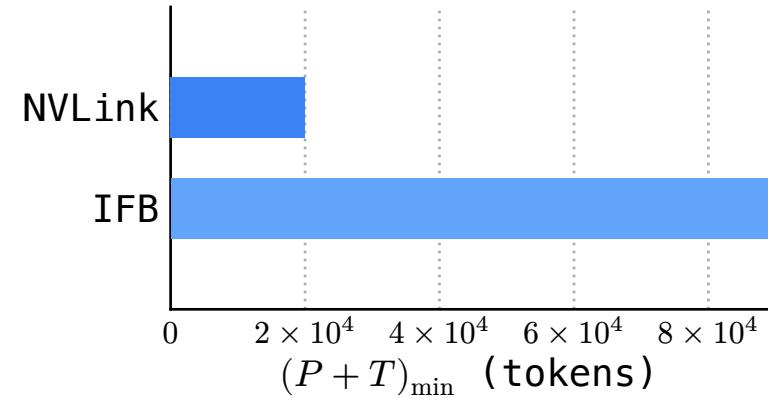
Overlap condition:

$$\frac{2T(P + T)D}{CN} \geq \frac{TDe}{\text{BW}}$$
$$P + T \geq \frac{N}{2} \frac{Ce}{\text{BW}}$$

Doesn't depend on  $N_H, N_{KV}$

# Q-Passing Roofline

Minimum  $(P + T)$  for communication overlap ( $8 \times$  B200):



Note:  $C \cdot e$  is the same for FP4 and FP8, so model doesn't matter!

Requires balanced KVs across ranks (round-robin during decode)

# This Paper Kinda Stinks

They have a bunch of math mistakes. Ex:

$$\text{Latency}(All2All) = (N - 1) \cdot \frac{(D + 1) \cdot T \cdot e}{BW}$$

# This Paper Kinda Stinks

They have a bunch of math mistakes. Ex:

$$\text{Latency}(All2All) = (N - 1) \cdot \frac{(D + 1) \cdot T \cdot e}{BW}$$

This is *so wrong it hurts.*

# This Paper Kinda Stinks

They have a bunch of math mistakes. Ex:

$$\text{Latency}(All2All) = (N - 1) \cdot \frac{(D + 1) \cdot T \cdot e}{BW}$$

This is *so wrong it hurts.*

They also don't actually solve their inequalities!

# This Paper Kinda Stinks

They have a bunch of math mistakes. Ex:

$$\text{Latency}(All2All) = (N - 1) \cdot \frac{(D + 1) \cdot T \cdot e}{BW}$$

This is *so wrong it hurts*.

They also don't actually solve their inequalities!

Let's do it right.

# All-to-All cost

All-to-All time in a ring is roughly<sup>1</sup>

$$T_{\text{all2all}} = \frac{TDe}{4 \text{ BW}}$$

All-to-All time in NVL72 is just  $\frac{TDe}{N \text{ BW}}$

Need to check if  $T_{\text{all2all}} < (T_{\text{kv,comm}} - T_{\text{kv,compute}})$

---

<sup>1</sup>See this derivation

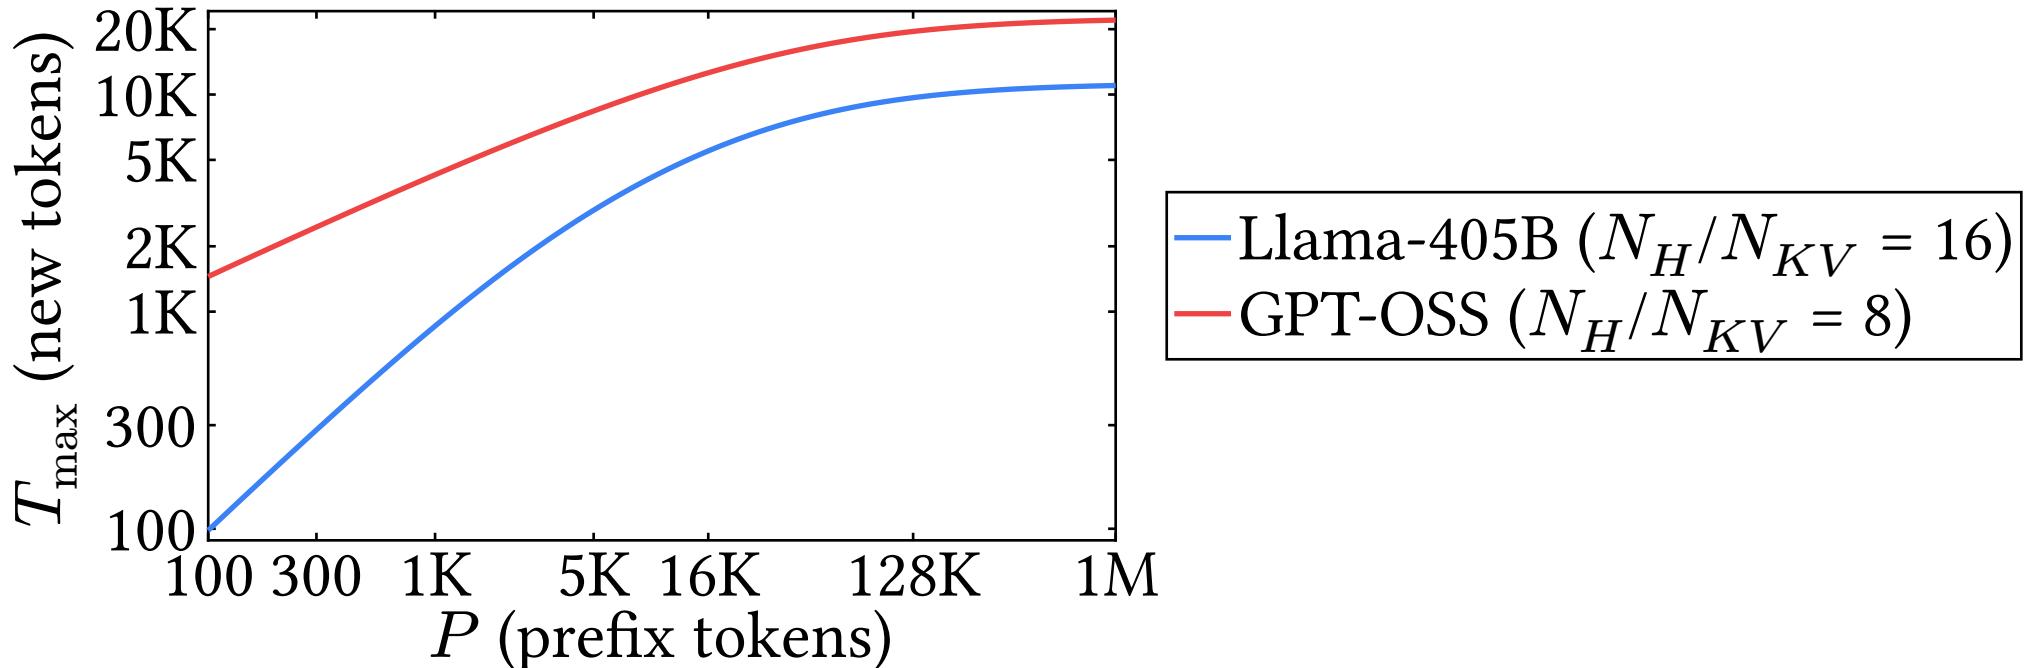
## All-to-All cost

$$\begin{aligned} T_{\text{kv,comm}} - T_{\text{kv,compute}} &= 2(P + T)D \frac{N_{KV}}{N_H} \frac{e}{\text{BW}} - \frac{2(P + T)TD}{NC} \\ &= 2(P + T)D \left( \frac{N_{KV}}{N_H} \frac{e}{\text{BW}} - \frac{T}{NC} \right) \\ \frac{TD e}{4 \text{ BW}} &< T_{\text{kv,comm}} - T_{\text{kv,compute}} \end{aligned}$$

Ends up being quadratic... hand it to `sympy`, solve for  $T$  and plot

# $T_{\max}$ vs Prefix Length

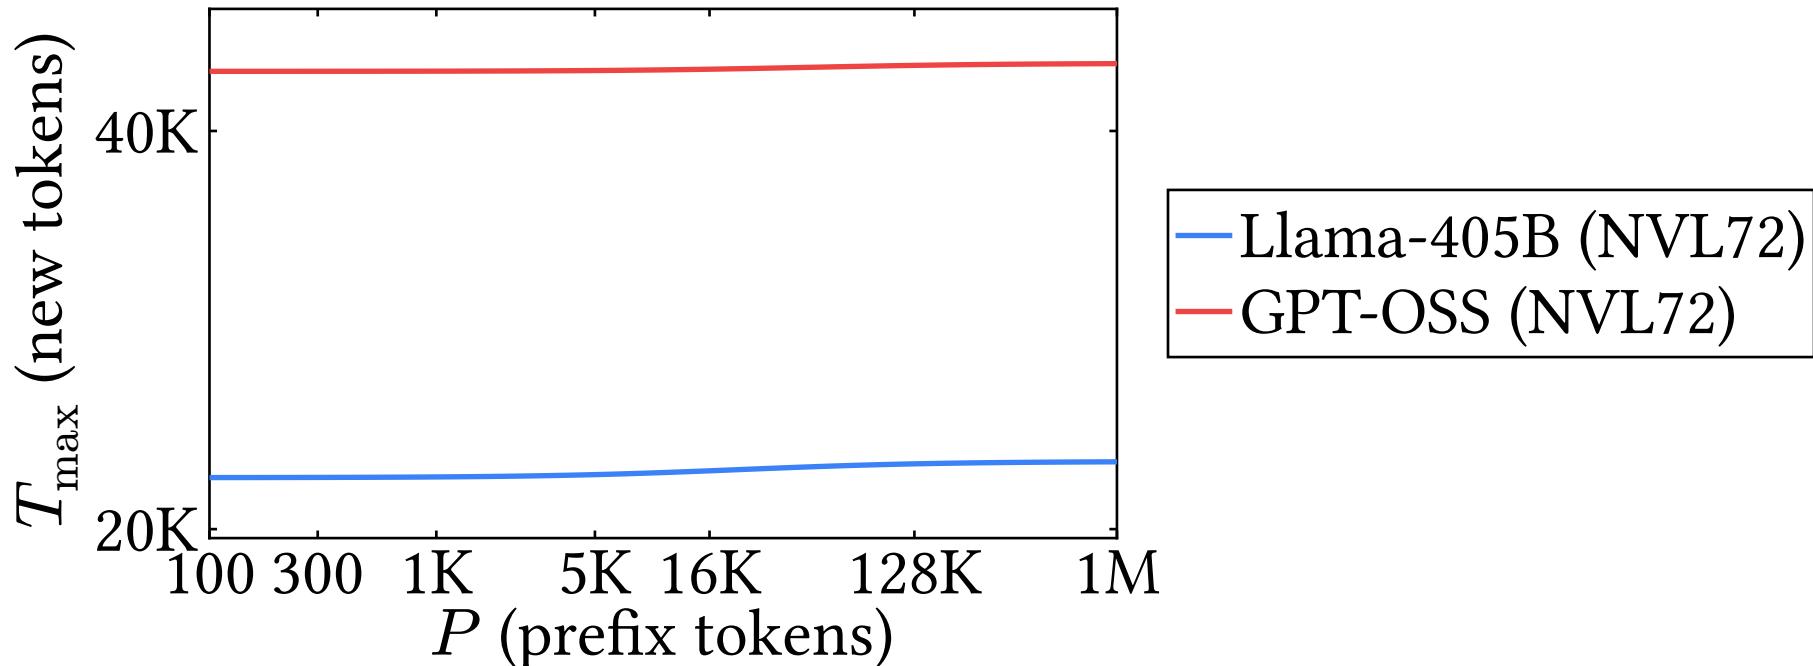
Maximum  $T$  where Pass-Q is faster than Pass-KV ( $8 \times$  B200 IFB):



Below  $T_{\max}$ , use All-to-All (Pass-Q). Above  $T_{\max}$ , use Pass-KV.

# What about NVL72?

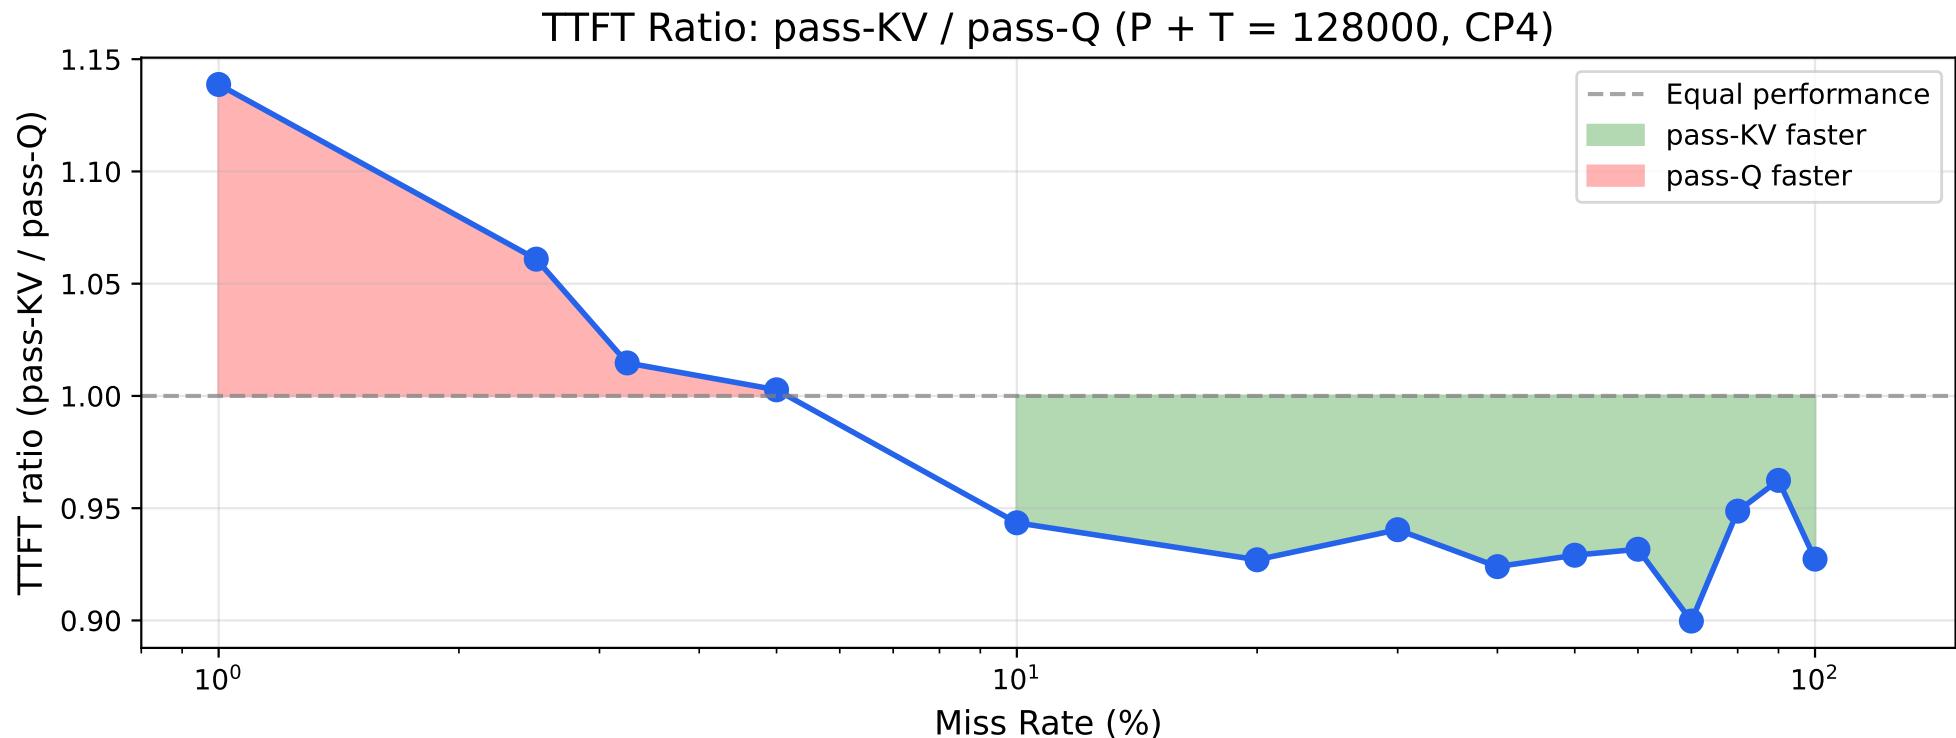
All-to-all cost is  $\frac{TDe}{N \cdot 4 \cdot \text{BW}}$  and BW is 900GB/s v.s. 200GB/s.



With NVL72's lower all-to-all cost, Pass-Q is optimal up to larger  $T$ .

But we only fail to hide communication for fairly small  $T$ .

# Meta's Hopper Performance

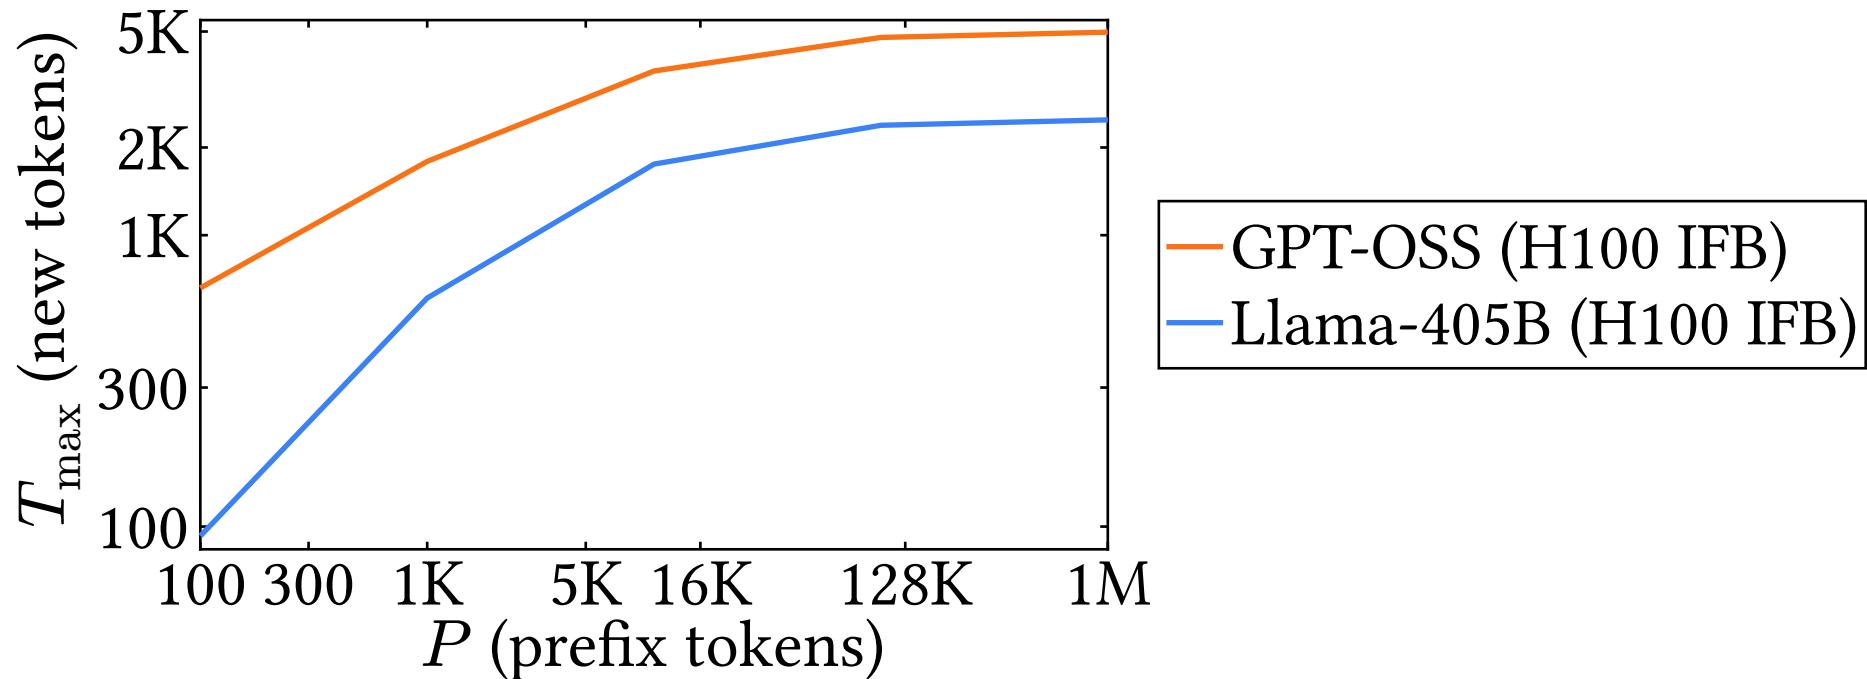


pass-Q is slightly better than pass-KV for very low miss rates

Does the math say the same?

# Validating Meta's Hopper Performance

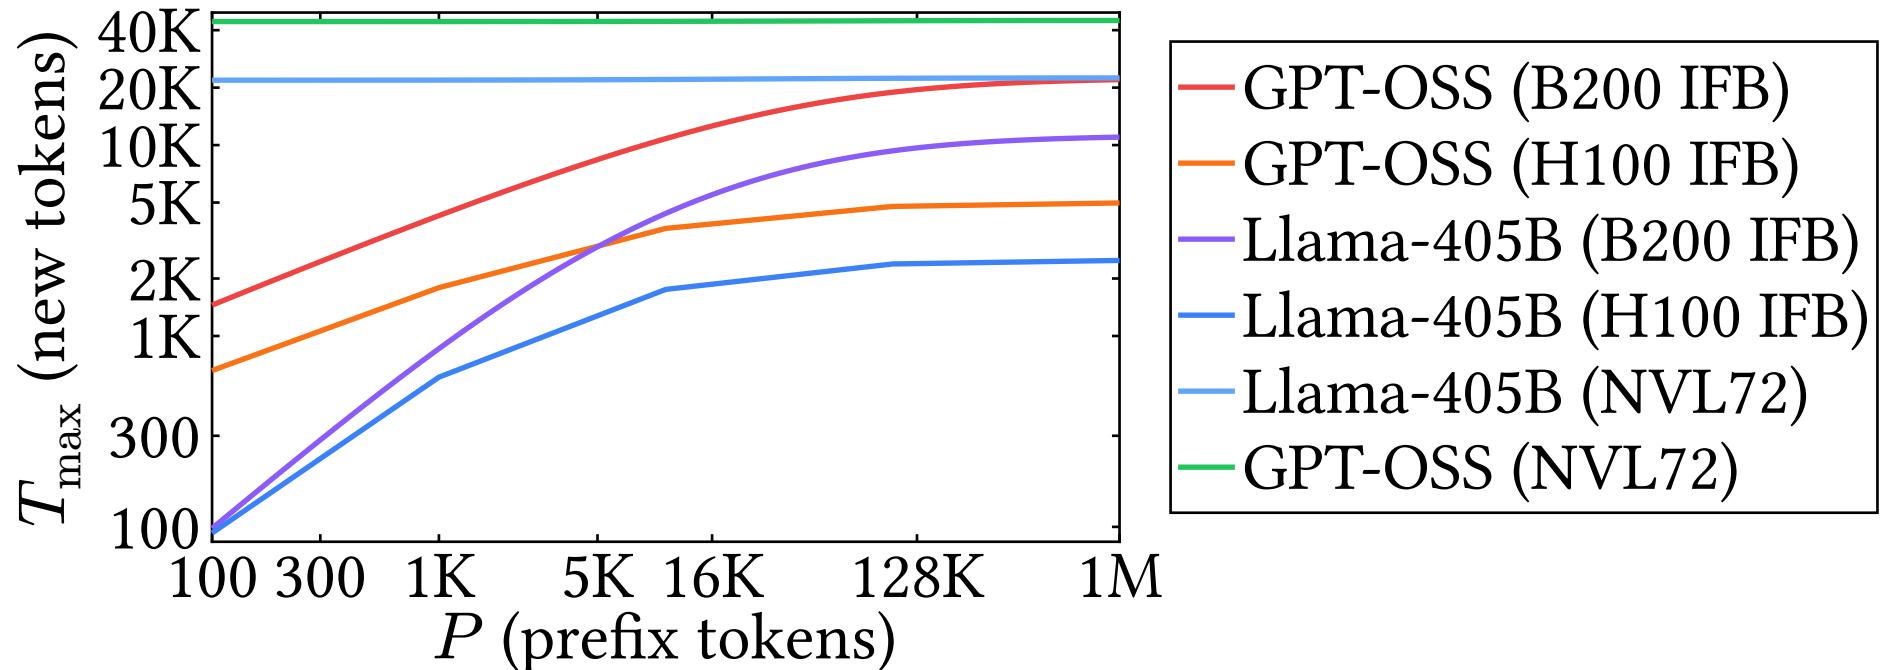
4 x H100 with 200GB/s (unidirectional) IFB



For H100 IFB at  $P=128K$ ,  $T_{\max} \approx 5K \approx 4\%$  miss rate! Empirically  $\approx 5\%$ !

# Putting it all together

All  $N = 8$



NVL72 pass-Q may be optimal, but  $T$  which isn't hideable is small.